

Ong Teck Wu
1001539
Lab 2

Files

meanThreadCode.c
medianThreadCode.c

How to run code

```
gcc -lpthreads meanThreadCode -o meanThreadCode.c  
./meanThreadCode input.txt 1 2 4 8 16 32 64 128 256 512 1024 2048
```

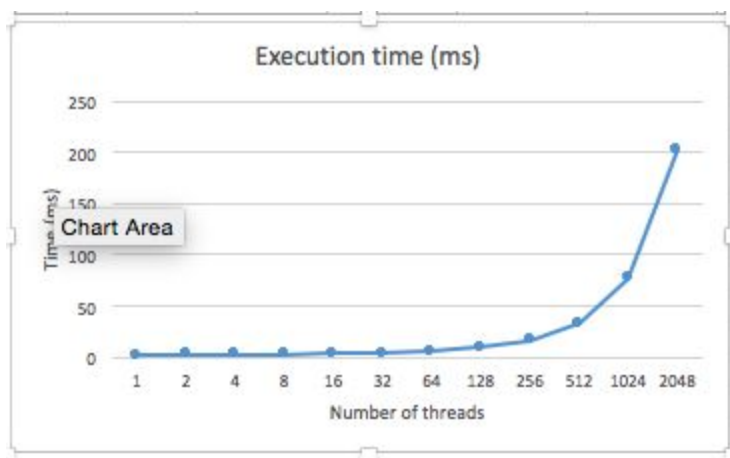
```
gcc -lpthreads medianThreadCode -o medianThreadCode.c  
./medianThreadCode input.txt 1 2 4 8 16 32 64 128 256 512 1024 2048
```

Set up

I performed the analysis by running the algorithm on each number of threads 5 times and getting the average run time.

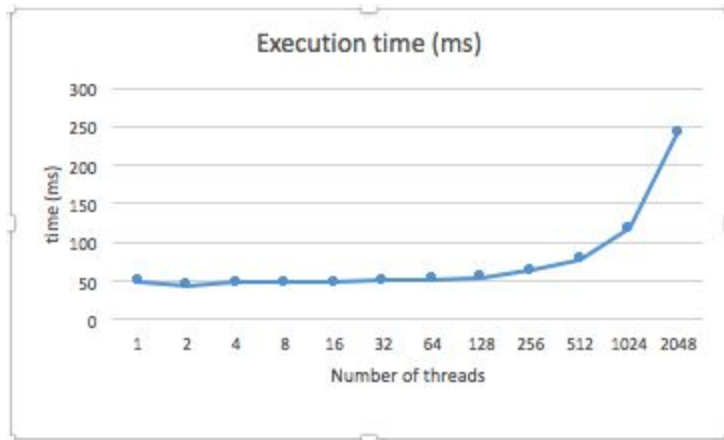
Mean

Number of threads	1	2	4	8	16	32	64	128	256	512	1024	2048
Execution time (ms)	2.13	2.18	2.72	2.85	3.25	4.05	5.49	9.28	16.75	33.17	76.77	202.37



Median

Number of threads	1	2	4	8	16	32	64	128	256	512	1024	2048
Execution time (ms)	49.3	44.2	48	47.9	48.2	50.4	52.4	54.9	64.1	78.3	118.	242.
	9	7			9	2	7	5	5	2	38	83



Observations

Mean: The execution time increases steadily but negligibly from 1-16 threads but starts to increase exponentially from 16-2048.

Median: The execution time increases steadily but negligibly from 1-128 threads but starts to increase exponentially from 128-2048. For median the minima is at 2 and the only after 16 did the runtime exceed that of 1. However, that is also negligible.

Analysis

For both mean and median, the runtime increases exponentially as a result of

- (1) the overhead from context switching,
- (2) having only two processors i.e. max two threads running at the same time,
- (3) and having to wait for the last worker thread to end before the next line of computation can start.

For such map-reduce computation operations, running more threads would only be ideal if there are more processors and an optimized maximum time-slice for processes to run in (to reduce redundant context switching overhead)

Problems I faced

For the “merge” function in medianThreadCode, due to the stack size having a limit of 2^{16} bytes, I had to use malloc to create temporary arrays for the naive implementation of the merge algorithm.